

Linear Regression Coding Assignment-2

Code

Hide

```
# Load essential libraries
library(ggplot2)
library(dplyr)
```

Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

filter, lag

The following objects are masked from 'package:base':

intersect, setdiff, setequal, union

Hide

```
library(HSAUR)
```

Warning: package 'HSAUR' was built under R version 4.3.2Loading required package: tools

Hide

```
library(ggcorrplot)
```

Warning: package 'ggcorrplot' was built under R version 4.3.2

Hide

```
# Load the heptathlon dataset
data(heptathlon)
str(heptathlon)
```

```
'data.frame': 25 obs. of 8 variables:
 $ hurdles : num 12.7 12.8 13.2 13.6 13.5 ...
 $ highjump: num 1.86 1.8 1.83 1.8 1.74 1.83 1.8 1.83 1.77 ...
 $ shot : num 15.8 16.2 14.2 15.2 14.8 ...
 $ run200m : num 22.6 23.6 23.1 23.9 23.9 ...
 $ longjump: num 7.27 6.71 6.68 6.25 6.32 6.33 6.37 6.47 6.11 6.28 ...
 $ javelin : num 45.7 42.6 44.5 42.8 47.5 ...
 $ run800m : num 129 126 124 132 128 ...
 $ score : int 7291 6897 6858 6540 6540 6411 6351 6297 6252 6252 ...
```

Hide

```
# Introduce a new column called sprint highlighting slow and fast sprinters
heptathlon = heptathlon %>% mutate(sprint = ifelse(run200m <= 25 & run800m <= 129, 'fast', 'slow'))
str(heptathlon)
```

```
'data.frame': 25 obs. of 9 variables:
 $ hurdles : num 12.7 12.8 13.2 13.6 13.5 ...
 $ highjump: num 1.86 1.8 1.83 1.8 1.74 1.83 1.8 1.83 1.77 ...
 $ shot : num 15.8 16.2 14.2 15.2 14.8 ...
 $ run200m : num 22.6 23.6 23.1 23.9 23.9 ...
 $ longjump: num 7.27 6.71 6.68 6.25 6.32 6.33 6.37 6.47 6.11 6.28 ...
 $ javelin : num 45.7 42.6 44.5 42.8 47.5 ...
 $ run800m : num 129 126 124 132 128 ...
 $ score : int 7291 6897 6858 6540 6540 6411 6351 6297 6252 6252 ...
 $ sprint : chr "fast" "fast" "fast" "slow" ...
```

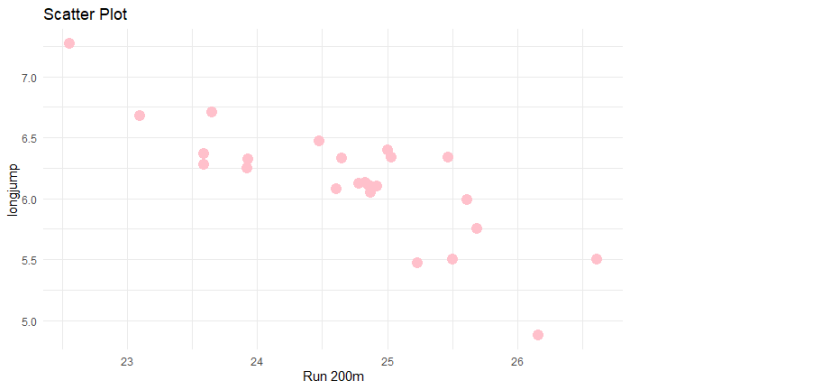
Hide

```
# Change sprint column to factor type
heptathlon['sprint'] = lapply(heptathlon['sprint'], as.factor)
str(heptathlon)
```

```
'data.frame': 25 obs. of 9 variables:
 $ hurdles : num 12.7 12.8 13.2 13.6 13.5 ...
 $ highjump: num 1.86 1.8 1.83 1.8 1.74 1.83 1.8 1.83 1.77 ...
 $ shot : num 15.8 16.2 14.2 15.2 14.8 ...
 $ run200m : num 22.6 23.6 23.1 23.9 23.9 ...
 $ longjump: num 7.27 6.71 6.68 6.25 6.32 6.33 6.37 6.47 6.11 6.28 ...
 $ javelin : num 45.7 42.6 44.5 42.8 47.5 ...
 $ run800m : num 129 126 124 132 128 ...
 $ score : int 7291 6897 6858 6540 6540 6411 6351 6297 6252 6252 ...
 $ sprint : Factor w/ 2 levels "fast","slow": 1 1 1 2 1 2 2 2 2 ...
```

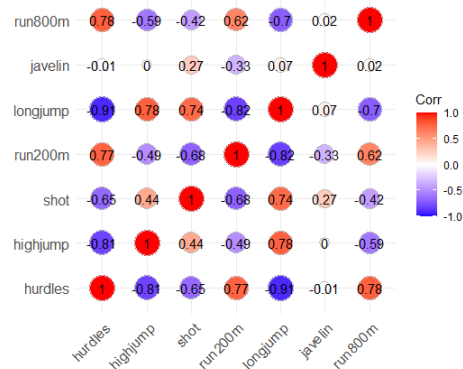
Hide

```
# Make a scatter plot between *run200m* (x-axis) and *longjump* (y-axis). What do you observe from this plot?
p = ggplot(heptathlon, aes(x=run200m,y=longjump))+
  geom_point(color='pink',size=4)+ labs(title = "Scatter Plot",x='Run 200m',y='longjump')+
  theme_minimal()
p
```



Hide

```
# Correlation between all pairs of continuous predictors (leave out sprint and the response variable score). What do you observe?
cor_matrix = cor(heptathlon %>% select(-c(sprint, score)))
ggcorrplot(cor_matrix, method = 'circle', lab = TRUE)
```




Hide

#Values close to 1 indicate a strong positive correlation, while values close to -1 indicate a strong negative correlation.

Hide

```
# Make a scatter plot between *run200m* (x-axis) and *longjump* (y-axis) now with the data points color-coded using *sprint*. What do you observe from this plot?
ggplot(heptathlon, aes(x = run200m, y = longjump, color = sprint)) +
  geom_point() +
  labs(title = "Scatter Plot: run200m vs longjump with Sprint Color Coding",
       x = "run200m", y = "longjump") + theme_minimal()
```

Scatter Plot: run200m vs longjump with Sprint Color Coding



Hide

# Calculate Pearson's correlation between \*run200m\* and \*longjump\*. What do you observe?

```
cor2 = cor(heptathlon['run200m'], heptathlon['longjump'], method = "pearson")
cor2
```

```
longjump
run200m -0.8172853
```

Hide

# How many levels does the categorical variable \*sprint\* have? What is the reference level?

```
contrasts(heptathlon$sprint)
```

```
slow
fast  0
slow  1
```

Hide

```
levels(heptathlon$sprint)
```

```
[1] "fast" "slow"
```

Hide

```
# Fit a linear model for approximating *score* as a function of *sprint*. Print the model's summary. How accurate is the model? How do the slow athletes' scores compare to the fast ones?
model = lm(data = heptathlon, score ~ sprint)
summary(model)
```

```
Call:
lm(formula = score ~ sprint, data = heptathlon)

Residuals:
    Min       1Q   Median       3Q      Max
-1347.4  -227.4   97.6   291.6   626.6

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  6799.4      200.3    33.939 < 2e-16 ***
sprintslow   -886.0      224.0    -3.956 0.000628 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 448 on 23 degrees of freedom
Multiple R-squared:  0.4049,    Adjusted R-squared:  0.379
F-statistic: 15.65 on 1 and 23 DF,  p-value: 0.0006282
```

Hide

```
mean_slow = mean(heptathlon[heptathlon$sprint == 'slow', 'score'])
mean_fast = mean(heptathlon[heptathlon$sprint == 'fast', 'score'])
mean_slow
```

```
[1] 5913.4
```

Hide

```
mean_fast
```

```
[1] 6799.4
```

Hide

```
mean_slow-mean_fast
```

```
[1] -886
```

Hide

```
# Fit a linear model for approximating *score* as a function of *shot* and *sprint*. Print the model's summary and answer the following questions:
```

```
# 1. Did the addition of the new predictor *shot* improve the model accuracy?
# 2. *True/false* (explain in one line): the model suggests that there is a possible linear relationship between an athlete's score and shotput performance.
# 3. For a 1 metre increase in shot put throw and with the same sprint performance, we can say with 95% confidence that the athlete's score will increase/decrease by an amount in the interval [?, ?].
model = lm(data = heptathlon, score ~ shot + sprint)
summary(model)
```

```
Call:
lm(formula = score ~ shot + sprint, data = heptathlon)
```

```
Residuals:
    Min       1Q   Median       3Q      Max
-1124.58  -164.40    35.93   207.34   496.35
```

```
Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  3080.0      883.0     3.488 0.002084 **
shot          249.7       58.4     4.275 0.000308 ***
sprintslow   -330.4      213.4    -1.548 0.135842
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 338.5 on 22 degrees of freedom
Multiple R-squared:  0.6749,    Adjusted R-squared:  0.6454
F-statistic: 22.84 on 2 and 22 DF,  p-value: 4.282e-06
```

Hide

```
# Using the model built above, extract the slope and intercept for estimating the *score* of *slow* and *fast* athletes.
```

```
# For slow athletes
intercept_slow = 3080.0
slope_slow = 249.7
```

```
# For fast athletes
intercept_fast = 2749.6
slope_fast = 249.7
```

Hide

```
# Complete the code below to build a linear model for approximating *score* as a function of *shot* and *sprint* using the training data. Predict the model performance by applying it to the test data.
# Split the data into 80% train and 20% test parts
set.seed(0)
train_ind = sample(1:nrow(heptathlon), size = 0.8*nrow(heptathlon))
```

```
hDataTrain = heptathlon[train_ind, ]
hDataTest = heptathlon[-train_ind, ]
```

```
# Build linear regression model
model = lm(score ~ shot + sprint, data = hDataTrain)
```

```
# Predict on the test data
predictions = predict(model, newdata = hDataTest)
```

```
# Print the true and predicted scores for the test data
print(cbind(TrueScore = hDataTest$score, PredictedScore = predictions))
```

```
      TrueScore PredictedScore
Behmer (GDR)    6858      6549.446
Greiner (USA)    6297      6279.790
Scheider (SWI)   6137      5592.081
Kytola (FIN)    5686      5613.656
Jeong-Mi (KOR)   5289      5389.814
```

Hide

```
# Calculate the model error (mean-squared error for test data)
mse = mean((hDataTest$score - predictions)^2)
print(paste("Mean Squared Error: ", mse))
```

```
[1] "Mean Squared Error: 81567.135660685"
```

Hide

```
# Fit a linear model for approximating *score* as a function of *shot*, *javelin*, and *sprint*. Print the model's summary and answer the following questions:
```

```
#1. Did the addition of the new predictor *javelin* improve the model accuracy?
#2. *True/false* (explain in one line): the model suggests that there is a possible linear relationship between an athlete's score and javelin performance.
#3. For a 1 metre increase in shot put throw and with the same javelin and sprint performance, we can say with 95% confidence that the athlete's score will increase/decrease by an amount in the interval [?, ?].
model = lm(score ~ shot + javelin + sprint, data = hDataTrain)
summary(model)
```

◀

▶

```
Call:
lm(formula = score ~ shot + javelin + sprint, data = hDataTrain)

Residuals:
    Min       1Q   Median       3Q      Max
-908.76 -113.96   8.78  204.83  449.76

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  3860.20    1521.11   2.538  0.02194 *
shot          277.65      71.55   3.880  0.00133 **
javelin       -28.24      26.35  -1.072  0.29966
sprintslow   -346.28     269.83  -1.283  0.21766
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 354.9 on 16 degrees of freedom
Multiple R-squared:  0.6808,    Adjusted R-squared:  0.621
F-statistic: 11.38 on 3 and 16 DF,  p-value: 0.0003043
```

Hide

```
# Fit a linear model for approximating *score* as a function of *highjump*, and *sprint*. Print the model's summary and answer the following questions:
# 1. How accurate is this model?
# 2. Considering a p-value of 10% as cutoff, are there any insignificant features?
model = lm(data = hDataTrain, score ~ highjump + sprint)
summary(model)
```

```
Call:
lm(formula = score ~ highjump + sprint, data = hDataTrain)

Residuals:
    Min       1Q   Median       3Q      Max
-481.77 -156.51  -10.06  122.85  476.28

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) -1893.2     1251.8  -1.512  0.148801
highjump     4801.1      689.2   6.966 2.28e-06 ***
sprintslow   -685.0      139.8  -4.902 0.000135 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 246.1 on 17 degrees of freedom
Multiple R-squared:  0.8369,    Adjusted R-squared:  0.8177
F-statistic: 43.62 on 2 and 17 DF,  p-value: 2.02e-07
```